

Programming Assignment 4

Little Search Engine

Posted Fri, Nov 4

Due Mon, Nov 21, 11:00 PM (WARNING!! NO GRACE PERIOD)

Extended deadline (with ONE time pass): Wed, Nov 23, 11:00 PM (NO GRACE PERIOD)

Worth 75 points = 7.5% of your course grade

In this assignment you will implement a simple search engine for text documents using hash tables.

You will work on this assignment individually. Read [DCS Academic Integrity Policy for Programming Assignments](#) - you are responsible for abiding by the policy. In particular, note that "All Violations of the Academic Integrity Policy will be reported by the instructor to the appropriate Dean".

You get ONE extension pass for the semester, no questions asked. There will be a total of 5 assignments this semester, and you may use this one time pass for any assignment. A separate Sakai assignment will be opened for extensions AFTER the deadline for the regular submission has passed.

IMPORTANT - READ THE FOLLOWING CAREFULLY!!!

- **Assignments emailed to the instructor or TAs will be ignored--they will NOT be accepted for grading. We will only grade submissions in Sakai.**
- **If your program does not compile, you will not get any credit.**

Most compilation errors occur for two reasons:

1. You are programming outside Eclipse, and you delete the "package" statement at the top of the file. If you do this, you are changing the program structure, and it will not compile when we test it.
2. You make some last minute changes, and submit without compiling.

To avoid these issues, (a) START EARLY, and give yourself plenty of time to work through the assignment, and (b) Submit a version well before the deadline so there is at least something in Sakai for us to grade. And you can keep submitting later versions (up to 10) - we will accept the LATEST version.

A separate Sakai assignment will be opened for extensions AFTER the deadline for the regular submission has passed.

Summary

You will implement a little search engine to do two things: (a) gather and index keywords that appear in a set of plain text documents, and (b) search for user-input keywords against the index and return a list of matching documents in which these keywords occur.

Implementation

Download the attached [lse_project.zip](#) file to your computer. DO NOT unzip it. Instead, follow the instructions on the Eclipse page under the section "Importing a Zipped Project into Eclipse" to get the entire project into your Eclipse workspace.

You will see a project called [Little Search Engine](#) with a single class, [search.LittleSearchEngine](#). This is where you will fill in your code, see details that follow.

The project also contains two sample text documents, [AliceCh1.txt](#), and [WowCh1.txt](#), directly under the project folder. You may use these samples to test your program. (Be sure to get other online text documents--or make your own--for more rigorous testing.) The names of these files are in a file called [docs.txt](#) that may be given as input to a driver that runs the search engine, which can in turn send this file name as the first argument to the [makeIndex](#) method in [LittleSearchEngine](#), to match with the [docsFile](#) parameter.

The project also has a [noisewords.txt](#) file that contains a list of "noise" words, one per line. Noise words are commonplace words (such as "the") that must be ignored by the search engine. You will use this file (and this file ONLY) to filter out noise words from the documents you read, when gathering keywords.

NOTE: You will need to write your own driver to test your implementation. This driver can be given a file that contains the names of all the documents, as well as the noise words file, as input. It can then set up a [LittleSearchEngine](#) object and call its methods as needed to test the implementation.

Following is the sequence of method calls that will be performed on a [LittleSearchEngine](#) object, to index and search keywords.

- [LittleSearchEngine\(\)](#) - Already implemented.

The constructor creates new (empty) [keywordsIndex](#) and [noiseWords](#) hash tables. The [keywordsIndex](#) hash table is the MASTER hash table, which indexes all keywords from all input documents. The [noiseWords](#) hash table stores all the noise words. Both of these are fields in the [LittleSearchEngine](#) class.

Every key in the [keywordsIndex](#) hash table is a keyword. The associated value for a keyword is an array list of (document,frequency) pairs for the documents in which the keyword occurs, *arranged in descending order of frequencies*. A (document,frequency) pair is held in an [Occurrence](#) object. The [Occurrence](#) class is defined in the [LittleSearchEngine.java](#) file, at the top. In an [Occurrence](#) object, the [document](#) field is the name of the document, which is basically the file name, e.g. [AliceCh1.txt](#).

- [void makeIndex\(String docsFile, String noiseWordsFile\)](#) - Already implemented.

Indexes all the keywords in all the input documents. See the method documentation and body in the [LittleSearchEngine.java](#) file for details.

If you want to index the given sample documents, the first parameter would be the file [docs.txt](#) and the second parameter would be the noise words file, [noisewords.txt](#)

After this method finishes executing, the full index of all keywords found in all input documents will be in the [keywordsIndex](#) hash table.

The [makeIndex](#) methods calls methods [loadKeyWords](#) and [mergeKeyWords](#), both of which you need to implement.

- [HashMap<String,Occurrence> loadKeyWords\(String docFile\)](#) - **You implement.**

This method creates a hash table for all keywords in a single given document. See the method documentation for details.

This method MUST call the `getKeyword` method, which you need to implement.

- `String getKeyword(String word)` - You implement.

Given an input word read from a document, it checks if the word is a keyword, and returns the keyword equivalent if it is.

See the method documentation for details. Here are some examples of input parameter word, and returned value.

Input Parameter	Returned value
distance.	distance (strip off period)
equi-distant	null (not all alphabetic characters)
Rabbit	rabbit (convert to lowercase)
Between	null (noise word)
we're	null (not all alphabetic characters)
World...	world (strip trailing periods)
World?!	world (strip trailing ? and !)
What,ever	null (not all alphabetic characters)

Note that if there is more than one trailing punctuation (as in the "World..." and "World?!" examples above), the method strips all of them. Also, the last example makes it clear that punctuation appearing anywhere but at the end is not stripped, and the word is rejected.

- `void mergeKeywords` - You implement.

Merges the keywords loaded from a single document (in method `loadKeywords`) into the global `keywordsIndex` hash table.

See the method documentation for details. This method MUST call the `insertLastOccurrence` method, which you need to implement.

- `ArrayList<Integer> insertLastOccurrence(ArrayList<Occurrence> occs)` - You implement.

See the method documentation for details. Note that this method uses binary search on frequency values to do the insertion. The return value is the sequence of mid points encountered during the search, using the regular (not lazy) binary search we covered in class. This return value is not used by the calling method-it is only going to be used for grading this method.

For example, suppose the list had the following frequency values (including the last one, which is to be inserted):

12	8	7	5	3	2	6
0	1	2	3	4	5	6

Then, the binary search (on the list *excluding* the last item) would encounter the following sequence of midpoint indexes:

2 4 3

Note that if a subarray has an even number of items, then the midpoint is the last item in the first half.

After inserting 6, the input list would be updated to this:

```
-----
12  8  7  6  5  3  2
-----
 0   1  2  3  4  5  6
```

and the sequence 2 4 3 would be returned.

If the new item is a duplicate of something that already exists, it doesn't matter if the new item is placed before or after the existing item.

Note that the items are in DESCENDING order, so the binary search would have to be done accordingly.

- `ArrayList<String> top5search(String kw1, String kw2)` - You implement.

This method computes the search result for the input "kw1 OR kw2", using the `keywordsIndex` hash table. The result is a list of NAMES of documents (limited to the top 5) in which either of the words "kw1" or "kw2" occurs, arranged in descending order of frequencies. See the method documentation in the code for additional details.

As an example, suppose the search is for "deep or world", in the given test documents, `AliceCh1.txt` (call it `A`) and `WowCh1.txt` (call it `W`). The word "deep" occurs twice in `A` and once in `W`, and the word "world" occurs once in `A` and 7 times in `W`:

```
deep:  (A,2), (W,1)
world: (W,7), (A,1)
```

The result of the search is:

```
WowCh1.txt, AliceCh1.txt
```

in that order. (Recall that the name of a document is the same as the name of the document file.)

NOTE:

- If a document occurs in both keywords' match list, consider the one with the higher frequency - do NOT add frequencies.
- Return AT MOST 5 non-duplicate entries. This means if there are more than 5 non-duplicate entries, then return the five top frequency entries, but if there are fewer than 5 non-duplicate entries, then return all of them.
- If a document in the first match list (for the first keyword) has the same frequency as a document in the second match list (for the second keyword), and both are candidates for inclusion in the output (they are not the same document), then pick the document in the first list before the document in the second list.

Implementation Rules

- Do NOT change the package name on the first line.
 - Do NOT add any import statements. With the existing imports, you may use any of the classes in `java.lang`, `java.io` and `java.util`.
 - Do NOT change the class `Occurrence` in ANY way.
 - Do NOT change the headers of any of the existing methods in `LittleSearchEngine` in ANY way.
 - Do NOT change the code in any of the implemented methods in ANY way.
 - Do NOT add any class fields in `LittleSearchEngine`.
 - You MAY add helper methods to `LittleSearchEngine`, but you must make them `private`.
-

Grading

Method	Points
<code>loadKeyWords</code>	15
<code>getKeyWord</code>	10
<code>mergeKeyWords</code>	15
<code>insertLastOccurrence</code>	15
<code>top5search</code>	20

When a method is graded, the correct versions of other methods will be used. Also, all data structures will be set to their correct (expected) states before a method is called.

Submission

Submit your `LittleSearchEngine.java` file.
